NVIDIA.

Jet-Nemotron: Efficient Language Model with Post Neural Architecture Search

Yuxian Gu, Qinghao Hu, Shang Yang, Haocheng Xi, Junyu Chen, Song Han, Han Cai

NVIDIA
https://github.com/NVlabs/Jet-Nemotron

Abstract: We present Jet-Nemotron, a new family of hybrid-architecture language models, which matches or exceeds the accuracy of leading full-attention models while significantly improving generation throughput. Jet-Nemotron is developed using Post Neural Architecture Search (PostNAS), a novel neural architecture exploration pipeline that enables efficient model design. Unlike prior approaches, PostNAS begins with a pre-trained full-attention model and freezes its MLP weights, allowing efficient exploration of attention block designs. The pipeline includes four key components: (1) learning optimal full-attention layer placement and elimination, (2) linear attention block selection, (3) designing new attention blocks, and (4) performing hardware-aware hyperparameter search. Our Jet-Nemotron-2B model achieves comparable or superior accuracy to Qwen3, Qwen2.5, Gemma3, and Llama3.2 across a comprehensive suite of benchmarks while delivering up to 53.6× generation throughput speedup and 6.1× prefilling speedup. It also achieves higher accuracy on MMLU and MMLU-Pro than recent advanced MoE full-attention models, such as DeepSeek-V3-Small and Moonlight, despite their larger scale with 15B total and 2.2B activated parameters.
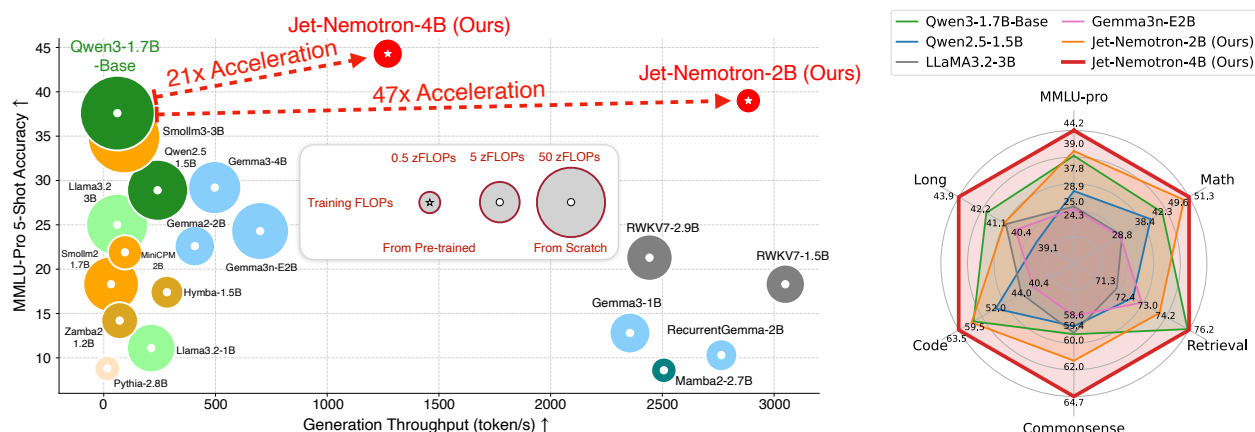
Figure 1 | **Comparison Between Jet-Nemotron and State-of-the-Art Efficient Language Models.** The generation throughput is measured on the NVIDIA H100 GPU under a context length of 64K tokens. Jet-Nemotron-2B delivers a higher accuracy than Qwen3-1.7B-Base on MMLU-Pro while achieving 47× higher generation throughput. Jet-Nemotron-4B, despite its larger model size, still achieves higher generation throughput than all full-attention models with less than 2B parameters.

1. Introduction

The rapid rise of Language Models (LMs) [1, 2, 3, 4, 5, 6, 7] marks a transformative era in artificial intelligence, with these models demonstrating exceptional accuracy across a broad range of tasks. However, their efficiency has become a significant concern due to the substantial computational and memory demands they impose. This issue is particularly pronounced in long-context generation and reasoning, where the self-attention mechanism [8] incurs a computational complexity of $O(n^2)$ and generates a large Key-Value (KV) cache[1].

To address this challenge, substantial efforts have been dedicated to designing more efficient LM architectures by developing attention mechanisms with reduced $O(n)$ complexity [9, 10, 11, 12, 13, 14]. In parallel, significant

---

[1]We refer to the standard $O(n^2)$ attention as full attention, and $O(n)$ attention as linear attention.

Figure 2 | **PostNAS Roadmap.** Our pipeline starts from a pre-trained full-attention model and keeps the MLP frozen. It then performs a coarse-to-fine search for efficient attention block designs, first determining the optimal placement of full-attention layers, then selecting the best linear attention block or using a new linear attention block, and finally searching for optimal architectural hyperparameters.

work has focused on constructing hybrid models that combine full and linear attention to strike a balance between accuracy and efficiency [15, 16, 17]. While these models offer improved efficiency over full-attention architectures, their accuracy still significantly falls behind state-of-the-art (SOTA) full-attention models, particularly on challenging benchmarks such as MMLU(-Pro) [18, 19], mathematical reasoning [20, 21, 22], retrieval [23, 24, 25], coding [26, 27, 28], and long-context tasks [29].

This paper introduces Jet-Nemotron, a new family of LMs that matches the accuracy of SOTA full-attention models while delivering exceptional efficiency. Figure 1 compares Jet-Nemotron with previous efficient LMs. Notably, Jet-Nemotron-2B achieves higher accuracy on MMLU-Pro than Qwen3-1.7B-Base [5], while offering 47× higher generation throughput on the NVIDIA H100 GPU under a context length of 64K.

Jet-Nemotron is built upon Post Neural Architecture Search (PostNAS), a novel neural architecture exploration pipeline (Figure 2) that enables the rapid design of efficient model architectures. Unlike the mainstream LM architecture design approaches, PostNAS begins with a pre-trained full-attention model, from which it inherits the Multi-Layer Perceptron (MLP) weights and keeps them frozen throughout the process. This strategy significantly reduces training costs while still allowing for comprehensive exploration of the attention block. The pipeline then proceeds through four key steps to systematically search for optimal attention block designs.

i) Full Attention Placement and Elimination. Retaining a few full-attention layers within the model [30] is essential for maintaining high accuracy on challenging tasks such as retrieval. However, the optimal placement of these layers remains unclear. In Section 2.2, we introduce a novel approach that automatically learns where to use full-attention layers by training a once-for-all super network [31] (Figure 4). The resulting learned placement significantly outperforms the commonly used uniform placement strategy in terms of accuracy on MMLU (Figure 5, right).

ii) Linear Attention Block Selection. After finalizing the placement of full-attention layers, we conduct an attention block search to identify the optimal linear attention block (Section 2.3). Thanks to the low training cost of our framework, we can systematically evaluate existing linear attention blocks in terms of accuracy across diverse tasks, training efficiency, and inference speed. Importantly, our approach eliminates the need to rely on small proxy tasks, such as training tiny LMs (e.g., 50M or 150M parameters), ensuring that the search results directly translate to improvements in final model accuracy. Moreover, as new linear attention blocks are out, our framework can rapidly evaluate them against prior designs and adopt them if they demonstrate promising results.

iii) New Attention Block Design. Our framework also facilitates the rapid design of new attention blocks. Adding convolutions is a widely used strategy to enhance the capacity of linear attention [32]. However, prior methods rely solely on static convolution kernels, lacking the ability to dynamically adapt convolution kernels' feature extraction patterns. In Section 2.4, we introduce a new linear attention block, JetBlock (Figure 2, #3). JetBlock uses a kernel generator to produce dynamic causal convolution kernels conditioned on the input, which are then applied to the value (V) tokens. Additionally, it removes redundant static convolutions on the query (Q) and key (K), streamlining the computation. Compared with prior linear attention blocks,

Figure 3 | **PostNAS Accuracy Improvement Breakdown.** By applying PostNAS to the baseline model, we achieve significant accuracy improvements across all benchmarks.

JetBlock shows improved accuracy with a small overhead (Table 1).

iv) Hardware-Aware Architecture Search. Last, in Section 2.5, we introduce a hardware-aware architecture search to identify optimal architectural hyperparameters. Traditionally, the number of parameters has been used as a proxy for LM efficiency. However, parameter count does not directly correlate with generation efficiency on actual hardware. Our hardware-aware search discovers architectural hyperparameters that deliver similar generation throughput, while using more parameters to achieve better accuracy (Table 2).

We evaluate Jet-Nemotron across a comprehensive suite of benchmarks, including MMLU(-Pro) [18, 19], commonsense reasoning [33, 34, 35, 36, 37, 38], mathematical reasoning [20, 21, 22, 39], retrieval [23, 24, 25], coding [26, 27, 28, 40], and long-context tasks [29]. Our Jet-Nemotron-2B model matches or surpasses SOTA full-attention models, such as Qwen2.5 [4], Qwen3 [5], Gemma3 [41, 42] and Llama3.2 [2], across all benchmarks, while achieving significantly higher generation throughput. Furthermore, the throughput gains are even more substantial in long-context settings (Figure 6). For example, with a 256K context length, Jet-Nemotron-2B delivers a 6.14× prefilling speedup and a 53.6× decoding speedup compared to Qwen3-1.7B-Base. We hope that our efficient LM family (Jet-Nemotron), our new linear attention block (JetBlock), and our architecture design pipeline (PostNAS) will benefit the community and accelerate the development and deployment of next-generation efficient LMs. We summarize our main contributions below:

- We introduce PostNAS, a novel model architecture exploration paradigm for language models. By reusing pre-trained LLMs, PostNAS reduces the cost and risk associated with LLM architecture exploration, enabling faster and more efficient innovation in the architecture design of LMs.
- We offer novel insights into the architecture design of efficient LMs, such as the task-specific importance of attention layers and the finding that KV cache size is a more critical factor than parameter count for generation throughput.
- We introduce a novel linear attention block, JetBlock, which integrates linear attention with dynamic convolution and hardware-aware architecture search. It consistently delivers significant accuracy improvements over previous linear attention blocks while maintaining comparable generation throughput.
- We introduce Jet-Nemotron, a novel hybrid-architecture LM family that achieves superior accuracy across a wide range of tasks and offers significantly higher generation throughput than prior SOTA full-attention models (e.g., Qwen2.5, Qwen3, Gemma3, and Llama3.2). With its strong accuracy and exceptional inference efficiency, Jet-Nemotron offers practical benefits for various applications requiring efficient LMs.

2. Method

2.1. PostNAS Motivation and Roadmap

Designing new language model architectures is challenging and risky due to the high cost of pre-training. Moreover, the significant gap in computational resources and training data makes it difficult for researchers outside of major organizations to match the accuracy of state-of-the-art full-attention models developed by large industry players [4, 41, 2]. This disparity hinders innovation in language model architecture design.

This paper proposes an alternative strategy for developing new language model architectures. Rather than pre-training models from scratch, we explore novel architectures by building on top of existing full-attention models. This approach dramatically reduces both training costs and data requirements.
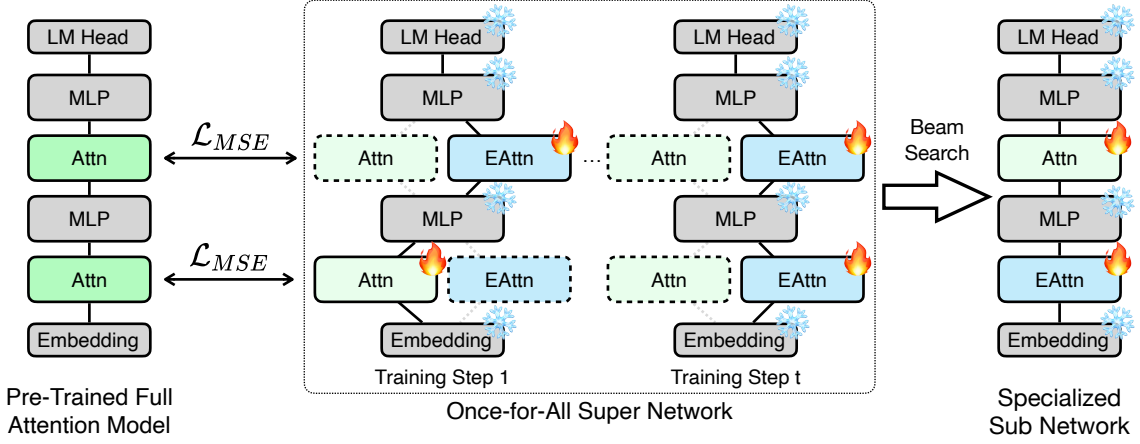
Figure 4 | **Learning to Place Full Attention with PostNAS.** We train a once-for-all super network and perform beam search to identify the optimal placement of full attention layers.


Figure 5 | **(a) Layer Placement Search Results on Qwen2.5-1.5B.** Each grid cell represents the search objective value of the corresponding attention layer; higher values indicate greater importance. **(b) Comparison Between PostNAS and Uniform Placement.**

While architectures designed within this framework may not yield optimal results when trained from scratch, we argue that they remain highly valuable. First, as demonstrated in Figure 1, they can deliver immediate gains in efficiency and accuracy over state-of-the-art full-attention models, translating to practical benefits such as improved services and reduced operational costs. Second, our framework serves as a rapid testbed for architectural innovation. If a new design fails to perform well in this setting, it will be unlikely to succeed in full pre-training [43]. This filtering mechanism helps researchers avoid wasting substantial computational resources on unpromising designs.

Figure 2 illustrates the roadmap of PostNAS. Starting from a pre-trained full-attention model, it freezes the MLP weights and explores attention block designs in a coarse-to-fine manner through four key steps: full attention placement and elimination (Section 2.2), linear attention block selection (Section 2.3), new attention block design (Section 2.4), and hardware-aware architecture search (Section 2.5). Figure 3 shows the accuracy improvement breakdown from these steps. We observe substantial accuracy improvements across all benchmarks: +5.3 on MMLU, +8.4 on math, +7.8 on retrieval, and +3.2 on commonsense reasoning.

2.2. Full Attention Placement and Elimination

Incorporating a few full-attention layers has become a common strategy for improving accuracy [30, 16, 44, 17]. The standard approach applies full attention uniformly across a fixed subset of layers, with the remaining layers using linear attention. However, this uniform strategy is suboptimal, especially in our setting, where we begin with a pre-trained full-attention model.

To address this, we propose an automatic method for efficiently determining the placement of full-attention layers. The overall approach is illustrated in Figure 4. We construct a once-for-all super network [45, 31] by augmenting the pre-trained full-attention model with alternative linear attention paths. During training, we

randomly sample an active path at each step, forming a subnetwork, which is trained using feature distillation loss [46, 47, 48].

After training, we perform beam search [49] to determine the optimal placement of full-attention layers under a given constraint (e.g., two full-attention layers). The search objective is task-dependent: for MMLU, we select the configuration with the lowest loss on the correct answer (i.e., maximizing $-loss$), while for mathematical and retrieval tasks, we choose the one with the highest accuracy. As shown in Figure 5(b), PostNAS significantly outperforms uniform placement in terms of accuracy.

Figure 5(a) presents the search results for Qwen2.5-1.5B. For each layer, we extract the corresponding subnetwork from the super network by configuring that layer as full attention while setting all remaining layers to linear attention. We evaluate the accuracy or loss of each subnetwork on a given task and visualize the results using a heatmap. Our analysis reveals three key findings:

> **Key Finding 1:** In the pre-trained full-attention model, not all attention layers contribute equally. For MMLU, only two layers exhibit critical importance, while for retrieval tasks, just two to three layers are particularly crucial.

> **Key Finding 2:** Different attention layers contribute to different capabilities. Layers that are critical for MMLU accuracy are not necessarily important for retrieval tasks.

> **Key Finding 3:** The pattern of attention importance becomes more intricate for complex tasks like mathematical reasoning. Fortunately, the combined set of top critical layers identified for MMLU and retrieval already encompasses most of the key layers needed for math.

In addition to these key findings, we observe that the search results remain consistent when using different linear attention operations. In our final experiments, we use GLA [11] in the once-for-all super network training for simplicity and slightly improved training throughput.

2.3. Linear Attention Block Selection

Building on the discovered full-attention layer placement, we conduct an attention block search to identify the most suitable linear attention block for our setup. In our experiments, we evaluate six SOTA linear attention blocks, including RWKV7 [10], RetNet [12], Mamba2 [50], GLA [11], Deltanet [51], and Gated DeltaNet [32].

After initial efficiency profiling, we observe that RWKV7 exhibits significantly lower training throughput compared to other linear attention blocks, possibly due to suboptimal kernel implementation. Consequently, we exclude it from our training experiments. The results, summarized in Table 1, indicate that Gated DeltaNet achieves the best overall accuracy among the evaluated linear attention blocks. This is attributed to the combination of two factors: (1) the Data-Dependent Gating Mechanism [52], which dynamically controls whether the model should focus more on the current token or the history state, and (2) the Delta Rule [53], which updates the history state with the information increment from the current token, to save the limited state memory. Therefore, we proceed with Gated DeltaNet in our experiments.

2.4. New Attention Block Design

We propose a new linear attention block, JetBlock, designed to enhance the model's expressive power by incorporating dynamic convolution [54, 55] into linear attention. Convolution has been shown to be essential for achieving strong accuracy in many linear attention blocks [32, 56]. However, prior works typically use static convolution kernels, which cannot adapt their feature extraction patterns based on the input.

To address this limitation, we introduce a kernel generator module that dynamically produces convolution kernels based on the input features. The overall structure is shown in Figure 2 (#3). This module shares the same input as the Q/K/V projection layer and begins with a linear reduction layer to improve efficiency, using a reduction ratio of 8. A SiLU activation function [57] is applied, followed by a final linear layer that outputs

Attention Block	Data-Depend Gating	Delta Rule	Throughput ↑		Accuracy ↑			
			Training	Inference	MMLU	Math	Retreival	Common.
RWKV7 [10]	✓	✓	123	2,542	-	-	-	-
RetNet [12]			269	2,535	53.6	29.9	63.7	58.1
Mamba2 [50]			273	3,220	51.5	26.0	68.9	57.5
GLA [11]	✓		265	3,079	55.8	31.2	66.6	58.5
Deltanet [51]		✓	254	2,955	48.9	27.4	67.9	56.6
Gated DeltaNet [32]	✓	✓	247	2,980	55.6	32.3	69.3	58.7
JetBlock	✓	✓	233	2,885	56.3	32.8	69.9	58.5
+ Hardware-Aware Search	✓	✓	227	2,883	**58.1**	**34.9**	**70.4**	**59.5**

Table 1 | **Accuracy and Efficiency of JetBlock**. JetBlock is designed through **Linear Attention Block Selection**, **New Attention Block Design**, and **Hardware-Aware Search** in PostNAS. It achieves higher accuracy than previous linear attention blocks while maintaining comparable training and inference efficiency.

d_K	d_V	n_{head}	Params (B)	Cache Size (MB)	Throughput (token/s) ↑	Retrieval Accuracy ↑	Math Accuracy ↑
256	288	4	1.62	154	2,969	67.6	31.3
192	384	4	1.64	154	2,961	69.3	32.3
128	576	4	1.70	154	2,979	69.5	32.5
256	144	8	1.66	154	2,986	68.3	32.1
192	192	8	1.70	154	2,970	70.6	32.8
128	288	8	1.74	154	2,971	69.6	33.2
128	192	12	1.78	154	2,959	68.8	32.9
96	256	12	1.84	154	2,955	69.6	34.8
64	384	12	1.98	154	2,952	70.1	34.2

Table 2 | **Detailed Results of Hardware-Aware Architecture Search.** The gray row is the original design [32], while the blue row shows the new design produced by our hardware-aware architecture search.

the convolution kernel weights. We adopt Gated DeltaNet for time-mixing, as it performs best compared with other designs as discussed in Section 2.3.

We apply the dynamic convolution kernels to the value (V) tokens, as applying them to the query (Q) or key (K) tokens offers little benefit. Furthermore, we find that static convolutions on Q and K can be removed with negligible impact on the final model accuracy once dynamic convolution is applied to V. We adopt this design in our final experiments for its slightly improved efficiency. Table 1 compares JetBlock with previous linear attention blocks. It provides better accuracy on math reasoning and retrieval tasks than Gated DeltaNet while maintaining similar efficiency.

2.5. Hardware-Aware Architecture Search

After finalizing the macro architecture, specifically the placement of full-attention layers, and selecting the linear attention block, we perform a hardware-aware architecture search to optimize core architectural hyperparameters, including key/value dimension and the number of attention heads.

Conventionally, parameter size is the primary efficiency metric used to guide model architecture design. However, this approach is suboptimal, as parameter count does not directly correlate with hardware efficiency. We address this limitation by using the generation throughput as a direct target for selecting architectural hyperparameters. We find that:

Type	Model	Params (B)	Cache Size (MB)	Throughput (token/s) ↑	MMLU Acc. ↑	MMLU-Pro Acc. ↑	BBH Acc. ↑
$O(n^2)$	Qwen2.5-1.5B [4]	1.5	1,792	241	59.5	28.9	44.1
	Qwen3-1.7B-Base [5]	1.7	7,168	61	60.3	37.8	54.2
	Llama3.2-3B [2]	3.0	7,168	60	54.9	25.0	47.1
	MiniCPM-2B-128K [58]	2.8	23,040	18	46.0	18.0	36.5
	MobileLLM-1.5B [59]	1.5	4,320	101	26.0	9.4	27.2
	Smollm2-1.7B [60]	1.7	12,288	32	48.5	18.3	35.1
	DeepSeek-V3-Small@1.3T [6]	2.2/15	-	-	53.3	-	-
	Moonlight@1.2T [61]	2.2/15	-	-	60.4	28.1	43.2
$O(n)$	Mamba-2.7B [50]	2.7	80	2,507	25.1	8.6	25.7
	RWKV7-1.5B [10]	1.5	24	3,050	41.0	13.4	15.9
	Rec.Gemma-2B [62]	2.0	16	2,355	28.6	12.8	33.3
Hybrid	Gemma3n-E2B [42]	2.0	768	701	53.9	24.3	45.1
	Hymba-1.5B [44]	1.5	240	180	49.7	17.4	29.8
	Zamba2-1.2B [16]	1.2	6,114	71	43.1	14.2	19.6
	Jet-Nemotron-2B	2.0	154	2,885	<u>60.8</u>	<u>39.0</u>	<u>58.3</u>
	Jet-Nemotron-4B	4.0	258	1,271	**65.2**	**44.2**	**65.0**

Table 3 | **Results on MMLU(-Pro) and BBH.** DeepSeek-V3-Small@1.3T and Moonlight@1.2T are MoE models with 2.2B activated and 15B total parameters, trained on 1.3T and 1.2T tokens, respectively.

> **Key Finding 4:** KV cache size is the most critical factor influencing long-context and long-generation throughput. When the KV cache size is constant, models with different parameter counts exhibit similar generation throughput (Table 2).

This is because the decoding stage is typically memory-bandwidth-bound rather than compute-bound. In long-context scenarios, the KV cache often consumes more memory than the model weights. Reducing its size decreases memory transfer time per decoding step and enables a larger batch size, thereby improving the generation throughput.

Based on Finding 4, we fix the KV cache size to match the original design and conduct a small-scale grid search over the key dimension, value dimension, and number of attention heads. Table 2 summarizes the results, where all variants use the same linear attention block (i.e., Gated DeltaNet) but have different configurations. The blue row represents our final design, while the gray row corresponds to the original design. Our final configuration achieves a generation throughput comparable to the original while incorporating more parameters and improving accuracy. From Table 1, we can see that our hardware-aware search in PostNAS boosts the JetBlock's accuracy, while maintaining training and inference throughput.

3. Experiments

3.1. Setup

Jet-Nemotron Model Family. We construct two versions of Jet-Nemotron with different parameter sizes: Jet-Nemotron-2B and Jet-Nemotron-4B. We use the Retrieval task to guide the placement of full attention layers and the MMLU task to guide the placement of sliding window attention (SWA) layers. Jet-Nemotron-2B is built upon Qwen2.5-1.5B [4], incorporating two full-attention layers (No. 15 and 20) for retrieval tasks and two sliding window attention (SWA) layers (No. 21 and 22) for multiple-choice tasks like MMLU. We find multiple-choice tasks mainly rely on the pattern-matching property of the softmax operation to route the knowledge of answers to their options. SWA effectively preserves the accuracy on such tasks. The remaining attention layers are replaced with JetBlock. Similarly, Jet-Nemotron-4B is based on Qwen2.5-3B and includes three full-attention layers (No. 18, 21, 33) and seven SWA layers (No. 6, 17, 20, 22, 23, 26, and 28). We summarize the final model architectures in Appendix A.1.

Type	Model	Throughput (token/s) ↑	Accuracy ↑					
			Avg.	GSM8K	MATH	MathQA	MMLU-Stem	GPQA
$O(n^2)$	Qwen2.5-1.5B [4]	241	38.4	62.4	13.1	34.4	52.7	29.4
	Qwen3-1.7B-Base [5]	61	42.3	62.8	16.7	46.0	50.8	27.9
	Llama3.2-3B [2]	60	28.8	25.8	8.6	34.2	45.3	30.1
	MiniCPM-2B-128K [58]	18	27.6	39.2	5.9	28.5	36.3	28.1
	Smollm2-1.7B [60]	32	28.9	30.3	9.2	33.7	41.3	30.1
$O(n)$	Mamba2-2.7B [50]	2,507	16.6	3.0	3.9	24.3	26.6	25.3
	RWKV7-1.5B [10]	2,669	18.3	5.6	0.8	27.2	34.9	23.0
	Rec.Gemma-2B [62]	2,355	20.8	13.9	7.6	25.3	28.5	28.6
Hybrid	Gemma3n-E2B [42]	701	28.3	24.9	10.1	31.1	45.7	31.8
	Hymba-1.5B [44]	180	23.1	17.9	0.8	28.0	40.9	27.9
	Zamba2-1.2B [16]	71	24.8	28.1	5.9	26.0	36.5	27.7
	Jet-Nemotron-2B	2,885	<u>49.6</u>	<u>76.2</u>	<u>23.3</u>	**53.8**	<u>62.7</u>	<u>32.1</u>
	Jet-Nemotron-4B	1,271	**51.3**	**78.7**	**25.2**	<u>52.5</u>	**65.6**	**34.6**

Table 4 | **Results on Math Tasks.**

Training Details. The training consists of two stages. In the first stage, we freeze the MLPs and train the model using a distillation loss. In the second stage, we perform full-model training. At the first stage, we use a combination of Nemotron-CC [63] and Redstone-QA [64] as our pre-training corpus and train Jet-Nemotron models for 50B tokens. This is also the setting in Section 2 where we perform PostNAS. At the second stage, we include more high-quality data from math [65] and coding [66, 67] domains into our data mixture. The models are then trained on 350B tokens. We summarize the experimental costs in Appendix A.2.

Evaluation Details. We evaluate Jet-Nemotron across mainstream benchmark settings: MMLU(-Pro) [18, 19], mathematical reasoning [18, 20, 21, 22], commonsense reasoning [33, 34, 35, 36, 37, 38], retrieval [23, 24, 25], coding [26, 27, 28, 40], and long-context tasks [29]. We compare our models against state-of-the-art full-attention models [2, 4, 5], linear attention models [10, 50], and hybrid models [41, 44]. We adopt 4-shot evaluation for GSM8K [22] and MATH [18] and 5-shot evaluation for GPQA [20] and MMLU-Pro [19]. We use the official implementation of EvalPlus [40] and CRUXEval [28] for coding tasks. For all other tasks, we use the zero-shot setting. All evaluations are based on LM-Evaluation-Harness [68].

Throughput Testbed. Our throughput evaluation was performed on a DGX H100 server, featuring 8 NVIDIA H100 GPUs, 2 Intel Xeon Platinum 8480C (112 cores) CPUs, and 2TB of RAM. For fair and consistent comparisons, we employ the latest available software versions. Specifically, our environment include Pytorch 2.7.0 and Triton 3.3.0. We implement the full-attention block with FlashAttention 2.7.4 [69] and linear attention blocks with Flash-Linear-Attention 0.2.1 [70]. Model inference is based on the Transformers 4.52.0 implementation [71]. The context length is 64K, except stated explicitly, and each model is tested on a single H100 GPU. We report the cache sizes for a 64K input context in Table 3. When testing the throughput, we adopt chunk-prefilling [72] and search for the chunk sizes to maximize the batch size for each model under the constraint of the GPU memory. In this way, we measure the highest achievable decoding throughput on the device. We list the batch sizes used for each model in Appendix A.3.

3.2. Main Results on Accuracy

Results on MMLU(-Pro) and BBH. Table 3 compares Jet-Nemotron with the most advanced efficient language models. Jet-Nemotron-2B achieves 47× higher throughput and has 47× smaller cache size than Qwen3-1.7B-Base, while delivering significantly better accuracy on MMLU, MMLU-Pro, and BBH. Jet-Nemotron-2B even outperforms recent MoE models like DeepSeek-V3-Small [6] and Moonlight [61] with larger activated parameters (2.2B) and much larger total parameters (15B). When scaled to 4B parameters, Jet-Nemotron-4B *still maintains a 21× throughput advantage against Qwen3-1.7B-Base.* Compared to other linear attention and hybrid models, Jet-Nemotron also achieves substantially higher accuracy.

Results on Math Tasks. Table 4 reports our results on math tasks. Jet-Nemotron-2B achieves an average

Model	Throughput (token/s) ↑	Accuracy ↑							
		Avg.	ARC-c	ARC-e	PIQA	Wino.	OBQA	BoolQ	TruthQA
Qwen2.5-1.5B [4]	241	59.4	45.4	71.2	75.8	63.8	40.2	72.8	46.6
Qwen3-1.7B-Base [5]	61	60.0	44.9	68.6	75.5	63.8	39.0	79.0	**48.8**
Llama3.2-3B [2]	60	59.9	46.6	72.0	<u>78.0</u>	<u>69.3</u>	40.4	73.9	39.3
MiniCPM-2B-128K [58]	18	57.6	41.0	69.4	75.5	63.8	40.6	74.7	38.3
Smollm2-1.7B [60]	32	59.7	47.0	73.3	77.7	66.2	<u>44.6</u>	72.5	36.7
Mamba2-2.7B [50]	2,507	57.2	42.1	70.5	76.1	62.7	41.4	71.5	36.1
RWKV7-1.5B [10]	3,050	59.7	46.3	75.7	77.4	67.6	**45.4**	70.5	34.7
Rec.Gemma-2B [62]	2,355	46.5	29.4	41.5	66.6	54.1	27.0	72.0	34.7
Gemma3n-E2B [42]	701	58.6	43.2	73.1	77.0	60.8	40.8	76.0	39.1
Hymba-1.5B [44]	180	61.2	46.9	<u>76.9</u>	77.7	66.2	41.0	80.8	39.0
Zamba2-1.2B [16]	71	58.0	44.4	66.8	77.4	65.6	42.8	70.8	38.5
Jet-Nemotron-2B	2,885	<u>62.0</u>	<u>48.6</u>	74.8	75.4	65.8	40.6	<u>81.2</u>	<u>47.8</u>
Jet-Nemotron-4B	1,271	**64.7**	**51.7**	**79.2**	**78.1**	**70.5**	43.6	**83.0**	46.6

Table 5 | **Results on Commonsense Tasks.**

Type	Model	Throughput (token/s) ↑	Accuracy ↑			
			Avg.	FDA	SWDE	Squad
$O(n^2)$	Qwen2.5-1.5B [4]	241	72.4	**82.8**	86.3	48.1
	Qwen3-1.7B-Base [5]	61	<u>76.1</u>	81.8	89.2	57.2
	Llama3.2-3B [2]	60	71.3	82.3	<u>89.6</u>	56.4
	MiniCPM-2B-128K [58]	18	72.6	72.3	86.4	**59.1**
	Smollm2-1.7B [60]	32	68.9	78.1	82.4	46.3
$O(n)$	Mamba2-2.7B [50]	2,507	57.0	51.7	74.3	45.1
	RWKV7-1.5B [10]	3,050	58.6	54.5	73.3	48.0
	Rec.Gemma-2.6B [62]	2,355	68.8	62.3	86.4	57.8
Hybrid	Gemma3n-E2B [73]	701	74.0	77.3	86.4	<u>58.2</u>
	Hymba-1.5B [44]	180	57.1	46.6	74.4	50.2
	Zamba2-1.2B [16]	71	66.4	73.8	80.7	44.8
	Jet-Nemotron-2B	2,885	74.2	80.4	85.7	56.6
	Jet-Nemotron-4B	1,271	**76.2**	<u>82.5</u>	**89.7**	56.4

Table 6 | **Results on Retrieval Tasks.**

accuracy of 49.6, surpassing Qwen3-1.7B-Base by 6.3 while being 47× faster. In contrast, prior linear attention and hybrid models are far behind Qwen3 on math tasks.

Results on Commonsense Reasoning Tasks. Table 5 summarizes the results on commonsense reasoning tasks. Qwen2.5 and Qwen3 are relatively weak in this domain. Nevertheless, Jet-Nemotron-2B, which uses Qwen2.5-1.5B as the starting point, still demonstrates strong results, achieving an average accuracy of 62.0, outperforming all baseline models.

Results on Retrieval Tasks. Table 6 presents the results on retrieval tasks. Jet-Nemotron-2B outperforms all baselines except Qwen3-1.7B-Base. When scaled to 4B, Jet-Nemotron-4B achieves the best average accuracy of 76.2, while still maintaining 21× speedup compared to Qwen3.

Results on Coding Tasks. Table 7 shows the results on coding tasks. Jet-Nemotron-2B performs comparably to Qwen3-1.7B-base and Jet-Nemotron-4B achieves a higher accuracy across all coding tasks while still delivering a large advantage on generation throughput against leading LMs like Qwen3-1.7B-Base.

Results on Long-Context Tasks. A common concern with linear and hybrid architectures is their accuracy on long-context tasks. In Table 8, we evaluate this on LongBench [29] up to a 64K context length. Our findings

Type	Model	Throughput (token/s) ↑	Accuracy ↑			
			Avg.	EvalPlus	CRUXEval-I-cot	CRUXEval-O-cot
$O(n^2)$	Qwen2.5-1.5B [4]	241	52.0	54.3	56.0	45.8
	Qwen3-1.7B-Base [5]	61	58.9	<u>62.8</u>	60.4	53.4
	Llama3.2-3B [2]	60	44.0	35.5	54.7	41.7
	MiniCPM-2B-128K [58]	18	34.2	40.7	29.9	31.9
	Smollm2-1.7B [60]	32	36.2	20.6	49.5	38.6
$O(n)$	Mamba2-2.7B [50]	2,507	14.0	12.0	9.3	20.7
	RWKV7-1.5B [10]	3,050	13.2	16.8	8.0	14.7
	Rec.Gemma-2.6B [62]	2,355	36.8	29.5	46.7	34.2
Hybrid	Gemma3n-E2B [73]	701	40.4	29.6	49.9	41.6
	Hymba-1.5B [44]	180	30.3	31.3	32.2	27.5
	Zamba2-1.2B [16]	71	20.1	12.7	21.1	26.4
	Jet-Nemotron-2B	2,885	<u>59.5</u>	60.8	<u>61.1</u>	<u>56.7</u>
	Jet-Nemotron-4B	1,271	**63.5**	**65.6**	**65.9**	**59.0**

Table 7 | **Results on Coding Tasks.**

Type	Model	Throughput (token/s) ↑	Accuracy ↑					
			Avg.	Few-Shot	Code	Sum.	Single-Doc	Multi-Doc
$O(n^2)$	Qwen2.5-1.5B [4]	241	39.1	63.9	57.2	26.3	28.3	19.9
	Qwen3-1.7B-Base [5]	61	<u>42.2</u>	<u>68.8</u>	48.1	**26.8**	**36.6**	**30.6**
	Llama3.2-3B [2]	60	39.9	65.2	58.0	24.3	27.6	24.6
	MiniCPM-2B-128K [58]	18	41.1	57.3	59.6	25.7	<u>33.4</u>	<u>29.6</u>
	Smollm2-1.7B [60]	32	21.3	38.9	28.6	16.0	13.2	9.8
$O(n)$	Mamba2-2.7B [50]	2,507	10.3	6.4	30.2	9.1	3.5	2.5
	RWKV7-1.5B [10]	3,050	14.2	10.6	21.1	18.1	12.8	8.7
	Rec.Gemma-2.6B [62]	2,355	24.1	31.8	56.7	12.9	9.2	9.6
Hybrid	Gemma2-2.6B [73]	388	22.9	28.7	52.0	12.6	13.9	7.3
	Gemma3n-E2B [73]	701	40.4	56.4	**67.2**	25.6	29.3	28.6
	Hymba-1.5B [44]	180	28.0	36.1	53.5	51.8	14.0	19.8
	Zamba2-1.2B [16]	71	9.2	10.0	20.1	10.2	3.8	1.7
	Jet-Nemotron-2B	2,885	41.1	68.7	58.1	26.0	30.8	21.9
	Jet-Nemotron-4B	1,271	**43.9**	**69.7**	<u>63.2</u>	<u>26.4</u>	32.5	27.5

Table 8 | **Results on Long-Context Tasks.**

show that Jet-Nemotron-2B, with two full-attention layers, achieves performance comparable to leading models like Qwen2.5-1.5B and Gemma3n-E2B, which feature considerably more such layers. Furthermore, our Jet-Nemotron-4B outperforms Qwen3-1.7B-Base while delivering a 21× speedup in generation throughput. These results substantially advance the frontier of the efficiency-accuracy trade-off in long-context tasks.

Summary. Jet-Nemotron-2B and Jet-Nemotron-4B perform comparably with or even better than the advanced full-attention model (Qwen3-1.7B-Base) across all six evaluation domains. With significantly fewer full-attention layers and smaller KV cache size, Jet-Nemotron-2B and Jet-Nemotron-4B deliver 47× and 21× higher generation throughput than Qwen3-1.7B-Base, respectively.

3.3. Efficiency Benchmark Results

Figure 6 shows the throughput comparison between Qwen3-1.7B-Base and Jet-Nemotron-2B across various context lengths. During the prefilling stage, Jet-Nemotron-2B is initially 1.14 and 1.15 times faster than Qwen3-1.7B-Base at shorter context lengths (4K and 8K). This can be further improved by designing a better optimized kernel implementation of the JetBlock. As the context length increases, the benefits of linear

Figure 6 | **Efficiency Comparison Across Different Context Lengths.** Jet-Nemotron-2B achieves up to a 6.14× speedup in prefilling and a 53.6× speedup in decoding compared to Qwen3-1.7B-Base.

attention become prominent, making Jet-Nemotron-2B achieve 6.14× speedup at a 256K context length.

During the decoding stage, Jet-Nemotron-2B consistently outperforms Qwen3-1.7B-Base by a large margin. Since Jet-Nemotron-2B includes 2 full-attention layers with 2 groups of key-value states, its theoretical maximum speedup is $14 \times 4 = 56$ times compared to Qwen3-1.7B-Base with 28 full-attention layers, where each layer contains 8 groups of key-value states. In our throughput testbed, Jet-Nemotron-2B achieves a 15.6× speedup at a 4K context length and up to a 53.6× speedup at a 256K context length, almost reaching the theoretical upper bound.

4. Related Work

Large language models (LLMs) are powerful but computationally intensive, motivating many works to build efficient model architectures for LLMs. One line of research focuses on designing efficient linear attention blocks [9, 10, 11, 12, 32, 50, 51, 74, 75, 76, 77, 78, 79, 80, 81, 82, 83, 82, 84] or log-linear attention [85] blocks to replace full attention blocks. Orthogonally, another line of research tries to combine full attention and linear attention to build hybrid models [13, 15, 16, 17, 44, 86, 87, 88]. These works typically focus on the pre-training setting, and their accuracy lags behind leading full-attention models. Recently, there are some efforts on linearizing LLMs with full attention replaced with linear attention [89, 90, 91, 92, 93, 94, 95, 96]. However, their model architecture are poorly optimized due to the large overhead of evaluating specific configuration, and thus their results are still inferior to SOTA full-attention models.

Our work is also related to neural architecture search (NAS) [45, 97, 98, 99, 100], a powerful technique for exploring the architectural design space and discovering novel model structures. In particular, hardware-aware neural architecture search [45] enables the development of specialized model architectures optimized for target hardware by training a once-for-all super-network [31], or leveraging layer-wise distillation [101, 102], etc. However, NAS has been rarely applied in the era of large language models (LLMs) due to the prohibitive cost of pretraining. Recent efforts have primarily focused on building flexible LLM architectures [103, 104], which can generate a range of subnetworks with varying depths and widths to accommodate different hardware platforms. Nevertheless, the architectural backbone of these subnetworks remains unchanged, relying entirely on full-attention layers.

5. Conclusion

We introduce Jet-Nemotron, a new family of hybrid-architecture language models that outperform state-of-the-art full-attention models — including Qwen3, Qwen2.5, Gemma3, and Llama3.2 — while delivering substantial efficiency gains, with up to 53.6× higher generation throughput on H100 GPUs (256K context length, maximum batch size). Jet-Nemotron is enabled by two key innovations: (1) Post Neural Architecture Search, a highly efficient post-training architecture adaptation pipeline applicable to any pre-trained Transformer model; and (2) the JetBlock, a novel linear attention block that significantly outperforms prior designs such as Mamba2, GLA, and Gated DeltaNet. Extensive empirical results show that Jet-Nemotron achieves major efficiency improvements without compromising accuracy across a broad range of benchmarks. Additionally, Jet-Nemotron significantly reduces the cost and risk associated with LLM architecture exploration, enabling faster and more efficient innovation in language model design.

References

[1] Tom Brown, Benjamin Mann, Nick Ryder, Melanie Subbiah, Jared D Kaplan, Prafulla Dhariwal, Arvind Neelakantan, Pranav Shyam, Girish Sastry, Amanda Askell, et al. Language models are few-shot learners. *Advances in neural information processing systems*, 33:1877–1901, 2020.

[2] Aaron Grattafiori, Abhimanyu Dubey, Abhinav Jauhri, Abhinav Pandey, Abhishek Kadian, Ahmad Al-Dahle, Aiesha Letman, Akhil Mathur, Alan Schelten, Alex Vaughan, et al. The llama 3 herd of models. *arXiv preprint arXiv:2407.21783*, 2024.

[3] Bo Adler, Niket Agarwal, Ashwath Aithal, Dong H Anh, Pallab Bhattacharya, Annika Brundyn, Jared Casper, Bryan Catanzaro, Sharon Clay, Jonathan Cohen, et al. Nemotron-4 340b technical report. *arXiv preprint arXiv:2406.11704*, 2024.

[4] An Yang, Baosong Yang, Beichen Zhang, Binyuan Hui, Bo Zheng, Bowen Yu, Chengyuan Li, Dayiheng Liu, Fei Huang, Haoran Wei, et al. Qwen2. 5 technical report. *arXiv preprint arXiv:2412.15115*, 2024.

[5] An Yang, Anfeng Li, Baosong Yang, Beichen Zhang, Binyuan Hui, Bo Zheng, Bowen Yu, Chang Gao, Chengen Huang, Chenxu Lv, et al. Qwen3 technical report. *arXiv preprint arXiv:2505.09388*, 2025.

[6] Aixin Liu, Bei Feng, Bing Xue, Bingxuan Wang, Bochao Wu, Chengda Lu, Chenggang Zhao, Chengqi Deng, Chenyu Zhang, Chong Ruan, et al. Deepseek-v3 technical report. *arXiv preprint arXiv:2412.19437*, 2024.

[7] Gemini Team, Rohan Anil, Sebastian Borgeaud, Jean-Baptiste Alayrac, Jiahui Yu, Radu Soricut, Johan Schalkwyk, Andrew M Dai, Anja Hauth, Katie Millican, et al. Gemini: a family of highly capable multimodal models. *arXiv preprint arXiv:2312.11805*, 2023.

[8] Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N Gomez, Łukasz Kaiser, and Illia Polosukhin. Attention is all you need. *Advances in neural information processing systems*, 30, 2017.

[9] Angelos Katharopoulos, Apoorv Vyas, Nikolaos Pappas, and François Fleuret. Transformers are rnns: Fast autoregressive transformers with linear attention. In *International conference on machine learning*, pages 5156–5165. PMLR, 2020.

[10] Bo Peng, Ruichong Zhang, Daniel Goldstein, Eric Alcaide, Xingjian Du, Haowen Hou, Jiaju Lin, Jiaxing Liu, Janna Lu, William Merrill, et al. Rwkv-7" goose" with expressive dynamic state evolution. *arXiv preprint arXiv:2503.14456*, 2025.

[11] Songlin Yang, Bailin Wang, Yikang Shen, Rameswar Panda, and Yoon Kim. Gated linear attention transformers with hardware-efficient training. *arXiv preprint arXiv:2312.06635*, 2023.

[12] Yutao Sun, Li Dong, Shaohan Huang, Shuming Ma, Yuqing Xia, Jilong Xue, Jianyong Wang, and Furu Wei. Retentive network: A successor to transformer for large language models. *arXiv preprint arXiv:2307.08621*, 2023.

[13] Aonian Li, Bangwei Gong, Bo Yang, Boji Shan, Chang Liu, Cheng Zhu, Chunhao Zhang, Congchao Guo, Da Chen, Dong Li, et al. Minimax-01: Scaling foundation models with lightning attention. *arXiv preprint arXiv:2501.08313*, 2025.

[14] Aili Chen, Aonian Li, Bangwei Gong, Binyang Jiang, Bo Fei, Bo Yang, Boji Shan, Changqing Yu, Chao Wang, Cheng Zhu, et al. Minimax-m1: Scaling test-time compute efficiently with lightning attention. *arXiv preprint arXiv:2506.13585*, 2025.

[15] Yutao Sun, Li Dong, Yi Zhu, Shaohan Huang, Wenhui Wang, Shuming Ma, Quanlu Zhang, Jianyong Wang, and Furu Wei. You only cache once: Decoder-decoder architectures for language models. *Advances in Neural Information Processing Systems*, 37:7339–7361, 2024.

[16] Paolo Glorioso, Quentin Anthony, Yury Tokpanov, Anna Golubeva, Vasudev Shyam, James Whittington, Jonathan Pilault, and Beren Millidge. The zamba2 suite: Technical report. *arXiv preprint arXiv:2411.15242*, 2024.

[17] Liliang Ren, Yang Liu, Yadong Lu, Yelong Shen, Chen Liang, and Weizhu Chen. Samba: Simple hybrid state space models for efficient unlimited context language modeling. *arXiv preprint arXiv:2406.07522*, 2024.

[18] Dan Hendrycks, Collin Burns, Steven Basart, Andy Zou, Mantas Mazeika, Dawn Song, and Jacob Steinhardt. Measuring massive multitask language understanding. *arXiv preprint arXiv:2009.03300*, 2020.

[19] Yubo Wang, Xueguang Ma, Ge Zhang, Yuansheng Ni, Abhranil Chandra, Shiguang Guo, Weiming Ren, Aaran Arulraj, Xuan He, Ziyan Jiang, et al. Mmlu-pro: A more robust and challenging multi-task language understanding benchmark. In *The Thirty-eight Conference on Neural Information Processing Systems Datasets and Benchmarks Track*, 2024.

[20] David Rein, Betty Li Hou, Asa Cooper Stickland, Jackson Petty, Richard Yuanzhe Pang, Julien Dirani, Julian Michael, and Samuel R Bowman. Gpqa: A graduate-level google-proof q&a benchmark. In *First Conference on Language Modeling*, 2024.

[21] Aida Amini, Saadia Gabriel, Peter Lin, Rik Koncel-Kedziorski, Yejin Choi, and Hannaneh Hajishirzi. Mathqa: Towards interpretable math word problem solving with operation-based formalisms. *arXiv preprint arXiv:1905.13319*, 2019.

[22] Karl Cobbe, Vineet Kosaraju, Mohammad Bavarian, Mark Chen, Heewoo Jun, Lukasz Kaiser, Matthias Plappert, Jerry Tworek, Jacob Hilton, Reiichiro Nakano, et al. Training verifiers to solve math word problems. *arXiv preprint arXiv:2110.14168*, 2021.

[23] Simran Arora, Brandon Yang, Sabri Eyuboglu, Avanika Narayan, Andrew Hojel, Immanuel Trummer, and Christopher Ré. Language models enable simple systems for generating structured views of heterogeneous data lakes. *Proc. VLDB Endow.*, 17(2):92–105, October 2023.

[24] Colin Lockard, Prashant Shiralkar, and Xin Luna Dong. OpenCeres: When open information extraction meets the semi-structured web. In Jill Burstein, Christy Doran, and Thamar Solorio, editors, *Proceedings of the 2019 Conference of the North American Chapter of the Association for Computational Linguistics: Human Language Technologies, Volume 1 (Long and Short Papers)*, June 2019.

[25] Pranav Rajpurkar, Jian Zhang, Konstantin Lopyrev, and Percy Liang. SQuAD: 100,000+ questions for machine comprehension of text. In *Proceedings of EMNLP*, 2016.

[26] Jacob Austin, Augustus Odena, Maxwell Nye, Maarten Bosma, Henryk Michalewski, David Dohan, Ellen Jiang, Carrie Cai, Michael Terry, Quoc Le, et al. Program synthesis with large language models. *arXiv preprint arXiv:2108.07732*, 2021.

[27] Mark Chen, Jerry Tworek, Heewoo Jun, Qiming Yuan, Henrique Ponde de Oliveira Pinto, Jared Kaplan, Harri Edwards, Yuri Burda, Nicholas Joseph, Greg Brockman, Alex Ray, Raul Puri, Gretchen Krueger, Michael Petrov, Heidy Khlaaf, Girish Sastry, Pamela Mishkin, Brooke Chan, Scott Gray, Nick Ryder, Mikhail Pavlov, Alethea Power, Lukasz Kaiser, Mohammad Bavarian, Clemens Winter, Philippe Tillet, Felipe Petroski Such, Dave Cummings, Matthias Plappert, Fotios Chantzis, Elizabeth Barnes, Ariel Herbert-Voss, William Hebgen Guss, Alex Nichol, Alex Paino, Nikolas Tezak, Jie Tang, Igor Babuschkin, Suchir Balaji, Shantanu Jain, William Saunders, Christopher Hesse, Andrew N. Carr, Jan Leike, Josh Achiam, Vedant Misra, Evan Morikawa, Alec Radford, Matthew Knight, Miles Brundage, Mira Murati, Katie Mayer, Peter Welinder, Bob McGrew, Dario Amodei, Sam McCandlish, Ilya Sutskever, and Wojciech Zaremba. Evaluating large language models trained on code, 2021.

[28] Alex Gu, Baptiste Roziere, Hugh James Leather, Armando Solar-Lezama, Gabriel Synnaeve, and Sida Wang. Cruxeval: A benchmark for code reasoning, understanding and execution. In *Forty-first International Conference on Machine Learning*, 2024.

[29] Yushi Bai, Xin Lv, Jiajie Zhang, Hongchang Lyu, Jiankai Tang, Zhidian Huang, Zhengxiao Du, Xiao Liu, Aohan Zeng, Lei Hou, Yuxiao Dong, Jie Tang, and Juanzi Li. LongBench: A bilingual, multitask benchmark for long context understanding. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 3119–3137, Bangkok, Thailand, August 2024. Association for Computational Linguistics.

[30] Roger Waleffe, Wonmin Byeon, Duncan Riach, Brandon Norick, Vijay Korthikanti, Tri Dao, Albert Gu, Ali Hatamizadeh, Sudhakar Singh, Deepak Narayanan, et al. An empirical study of mamba-based language models. *arXiv preprint arXiv:2406.07887*, 2024.

[31] Han Cai, Chuang Gan, Tianzhe Wang, Zhekai Zhang, and Song Han. Once-for-all: Train one network and specialize it for efficient deployment. *arXiv preprint arXiv:1908.09791*, 2019.

[32] Songlin Yang, Jan Kautz, and Ali Hatamizadeh. Gated delta networks: Improving mamba2 with delta rule. *arXiv preprint arXiv:2412.06464*, 2024.

[33] Christopher Clark, Kenton Lee, Ming-Wei Chang, Tom Kwiatkowski, Michael Collins, and Kristina Toutanova. BoolQ: Exploring the surprising difficulty of natural yes/no questions. In *Proceedings of NAACL-HLT*, 2019.

[34] Peter Clark, Isaac Cowhey, Oren Etzioni, Tushar Khot, Ashish Sabharwal, Carissa Schoenick, and Oyvind Tafjord. Think you have solved question answering? try arc, the ai2 reasoning challenge. *arXiv:1803.05457v1*, 2018.

[35] Yonatan Bisk, Rowan Zellers, Jianfeng Gao, Yejin Choi, et al. Piqa: Reasoning about physical commonsense in natural language. In *Proceedings of AAAI*, 2020.

[36] Hector Levesque, Ernest Davis, and Leora Morgenstern. The winograd schema challenge. In *Proceedings of KR*, 2012.

[37] Stephanie Lin, Jacob Hilton, and Owain Evans. TruthfulQA: Measuring how models mimic human falsehoods. In *Proceedings of the 60th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, 2022.

[38] Todor Mihaylov, Peter Clark, Tushar Khot, and Ashish Sabharwal. Can a suit of armor conduct electricity? a new dataset for open book question answering. In *Proceedings of EMNLP*, 2018.

[39] Dan Hendrycks, Collin Burns, Saurav Kadavath, Akul Arora, Steven Basart, Eric Tang, Dawn Song, and Jacob Steinhardt. Measuring mathematical problem solving with the math dataset. *arXiv preprint arXiv:2103.03874*, 2021.

[40] Jiawei Liu, Chunqiu Steven Xia, Yuyao Wang, and Lingming Zhang. Is your code generated by chatGPT really correct? rigorous evaluation of large language models for code generation. In *Thirty-seventh Conference on Neural Information Processing Systems*, 2023.

[41] Gemma Team, Aishwarya Kamath, Johan Ferret, Shreya Pathak, Nino Vieillard, Ramona Merhej, Sarah Perrin, Tatiana Matejovicova, Alexandre Ramé, Morgane Rivière, et al. Gemma 3 technical report. *arXiv preprint arXiv:2503.19786*, 2025.

[42] Fnu Devvrit, Sneha Kudugunta, Aditya Kusupati, Tim Dettmers, Kaifeng Chen, Inderjit Dhillon, Yulia Tsvetkov, Hannaneh Hajishirzi, Sham Kakade, Ali Farhadi, and Prateek Jain. Matformer: Nested transformer for elastic inference. In *Workshop on Advancing Neural Network Training: Computational Efficiency, Scalability, and Resource Optimization (WANT@NeurIPS 2023)*, 2023.

[43] Barret Zoph, Vijay Vasudevan, Jonathon Shlens, and Quoc V Le. Learning transferable architectures for scalable image recognition. In *Proceedings of the IEEE conference on computer vision and pattern recognition*, pages 8697–8710, 2018.

[44] Xin Dong, Yonggan Fu, Shizhe Diao, Wonmin Byeon, Zijia Chen, Ameya Sunil Mahabaleshwarkar, Shih-Yang Liu, Matthijs Van Keirsbilck, Min-Hung Chen, Yoshi Suhara, et al. Hymba: A hybrid-head architecture for small language models. *arXiv preprint arXiv:2411.13676*, 2024.

[45] Han Cai, Ligeng Zhu, and Song Han. Proxylessnas: Direct neural architecture search on target task and hardware. *arXiv preprint arXiv:1812.00332*, 2018.

[46] Geoffrey Hinton, Oriol Vinyals, and Jeff Dean. Distilling the knowledge in a neural network. *arXiv preprint arXiv:1503.02531*, 2015.

[47] Yuxian Gu, Li Dong, Furu Wei, and Minlie Huang. MiniLLM: Knowledge distillation of large language models. In *Proceedings of ICLR*, 2024.

[48] Wenhui Wang, Furu Wei, Li Dong, Hangbo Bao, Nan Yang, and Ming Zhou. Minilm: Deep self-attention distillation for task-agnostic compression of pre-trained transformers. *Advances in neural information processing systems*, 33:5776–5788, 2020.

[49] Alex Graves. Sequence transduction with recurrent neural networks. In *Proceedings of the Workshop on Representation Learning (ICML2012)*, 2012.

[50] Tri Dao and Albert Gu. Transformers are ssms: Generalized models and efficient algorithms through structured state space duality. *arXiv preprint arXiv:2405.21060*, 2024.

[51] Songlin Yang, Bailin Wang, Yu Zhang, Yikang Shen, and Yoon Kim. Parallelizing linear transformers with the delta rule over sequence length. *arXiv preprint arXiv:2406.06484*, 2024.

[52] Jos Van Der Westhuizen and Joan Lasenby. The unreasonable effectiveness of the forget gate. *arXiv preprint arXiv:1804.04849*, 2018.

[53] DL Prados and SC Kak. Neural network capacity using delta rule. *Electronics Letters*, 25(3):197–199, 1989.

[54] Felix Wu, Angela Fan, Alexei Baevski, Yann N Dauphin, and Michael Auli. Pay less attention with lightweight and dynamic convolutions. *arXiv preprint arXiv:1901.10430*, 2019.

[55] Han Cai, Muyang Li, Qinsheng Zhang, Ming-Yu Liu, and Song Han. Condition-aware neural network for controlled image generation. In *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, pages 7194–7203, 2024.

[56] Han Cai, Junyan Li, Muyan Hu, Chuang Gan, and Song Han. Efficientvit: Lightweight multi-scale attention for high-resolution dense prediction. In *Proceedings of the IEEE/CVF international conference on computer vision*, pages 17302–17313, 2023.

[57] Stefan Elfwing, Eiji Uchibe, and Kenji Doya. Sigmoid-weighted linear units for neural network function approximation in reinforcement learning. *Neural networks*, 107:3–11, 2018.

[58] Shengding Hu, Yuge Tu, Xu Han, Chaoqun He, Ganqu Cui, Xiang Long, Zhi Zheng, Yewei Fang, Yuxiang Huang, Weilin Zhao, et al. Minicpm: Unveiling the potential of small language models with scalable training strategies. *arXiv preprint arXiv:2404.06395*, 2024.

[59] Zechun Liu, Changsheng Zhao, Forrest Iandola, Chen Lai, Yuandong Tian, Igor Fedorov, Yunyang Xiong, Ernie Chang, Yangyang Shi, Raghuraman Krishnamoorthi, et al. Mobilellm: Optimizing sub-billion parameter language models for on-device use cases. In *Forty-first International Conference on Machine Learning*, 2024.

[60] Loubna Ben Allal, Anton Lozhkov, Elie Bakouch, Gabriel Martín Blázquez, Guilherme Penedo, Lewis Tunstall, Andrés Marafioti, Hynek Kydlíček, Agustín Piqueres Lajarín, Vaibhav Srivastav, et al. Smollm2: When smol goes big–data-centric training of a small language model. *arXiv preprint arXiv:2502.02737*, 2025.

[61] Jingyuan Liu, Jianlin Su, Xingcheng Yao, Zhejun Jiang, Guokun Lai, Yulun Du, Yidao Qin, Weixin Xu, Enzhe Lu, Junjie Yan, et al. Muon is scalable for llm training. *arXiv preprint arXiv:2502.16982*, 2025.

[62] Aleksandar Botev, Soham De, Samuel L Smith, Anushan Fernando, George-Cristian Muraru, Ruba Haroun, Leonard Berrada, Razvan Pascanu, Pier Giuseppe Sessa, Robert Dadashi, et al. Recurrentgemma: Moving past transformers for efficient open language models. *arXiv preprint arXiv:2404.07839*, 2024.

[63] Dan Su, Kezhi Kong, Ying Lin, Joseph Jennings, Brandon Norick, Markus Kliegl, Mostofa Patwary, Mohammad Shoeybi, and Bryan Catanzaro. Nemotron-cc: Transforming common crawl into a refined long-horizon pretraining dataset. *arXiv preprint arXiv:2412.02595*, 2024.

[64] Yaoyao Chang, Lei Cui, Li Dong, Shaohan Huang, Yangyu Huang, Yupan Huang, Scarlett Li, Tengchao Lv, Shuming Ma, Qinzheng Sun, et al. Redstone: Curating general, code, math, and qa data for large language models. *arXiv preprint arXiv:2412.03398*, 2024.

[65] Fan Zhou, Zengzhi Wang, Nikhil Ranjan, Zhoujun Cheng, Liping Tang, Guowei He, Zhengzhong Liu, and Eric P Xing. Megamath: Pushing the limits of open math corpora. *arXiv preprint arXiv:2504.02807*, 2025.

[66] Loubna Ben Allal, Anton Lozhkov, Guilherme Penedo, Thomas Wolf, and Leandro von Werra. Smollm-corpus, July 2024.

[67] Kazuki Fujii, Yukito Tajima, Sakae Mizuki, Hinari Shimada, Taihei Shiotani, Koshiro Saito, Masanari Ohi, Masaki Kawamura, Taishi Nakamura, Takumi Okamoto, et al. Rewriting pre-training data boosts llm performance in math and code. *arXiv preprint arXiv:2505.02881*, 2025.

[68] Leo Gao, Jonathan Tow, Baber Abbasi, Stella Biderman, Sid Black, Anthony DiPofi, Charles Foster, Laurence Golding, Jeffrey Hsu, Alain Le Noac'h, Haonan Li, Kyle McDonell, Niklas Muennighoff, Chris Ociepa, Jason Phang, Laria Reynolds, Hailey Schoelkopf, Aviya Skowron, Lintang Sutawika, Eric Tang, Anish Thite, Ben Wang, Kevin Wang, and Andy Zou. The language model evaluation harness, 07 2024.

[69] Tri Dao. Flashattention-2: Faster attention with better parallelism and work partitioning. In *The Twelfth International Conference on Learning Representations*, 2024.

[70] Songlin Yang and Yu Zhang. Fla: A triton-based library for hardware-efficient implementations of linear attention mechanism, January 2024.

[71] Thomas Wolf, Lysandre Debut, Victor Sanh, Julien Chaumond, Clement Delangue, Anthony Moi, et al. Transformers: State-of-the-art natural language processing. In *Proceedings of the 2020 Conference on Empirical Methods in Natural Language Processing: System Demonstrations*, 2020.

[72] Amey Agrawal, Nitin Kedia, Ashish Panwar, Jayashree Mohan, Nipun Kwatra, Bhargav Gulavani, Alexey Tumanov, and Ramachandran Ramjee. Taming Throughput-Latency tradeoff in LLM inference with Sarathi-Serve. In *18th USENIX Symposium on Operating Systems Design and Implementation (OSDI 24)*, pages 117–134, Santa Clara, CA, July 2024. USENIX Association.

[73] Gemma Team, Morgane Riviere, Shreya Pathak, Pier Giuseppe Sessa, Cassidy Hardin, Surya Bhupatiraju, Léonard Hussenot, Thomas Mesnard, Bobak Shahriari, Alexandre Ramé, et al. Gemma 2: Improving open language models at a practical size. *arXiv preprint arXiv:2408.00118*, 2024.

[74] Zhen Qin, Songlin Yang, and Yiran Zhong. Hierarchically gated recurrent neural network for sequence modeling. *Advances in Neural Information Processing Systems*, 36:33202–33221, 2023.

[75] Albert Gu, Karan Goel, and Christopher Ré. Efficiently modeling long sequences with structured state spaces. *arXiv preprint arXiv:2111.00396*, 2021.

[76] Weizhe Hua, Zihang Dai, Hanxiao Liu, and Quoc Le. Transformer quality in linear time. In *International conference on machine learning*, pages 9099–9117. PMLR, 2022.

[77] Zhen Qin, Yuxin Mao, Xuyang Shen, Dong Li, Jing Zhang, Yuchao Dai, and Yiran Zhong. You only scan once: Efficient multi-dimension sequential modeling with lightnet. *arXiv preprint arXiv:2405.21022*, 2024.

[78] Zhen Qin, Weixuan Sun, Hui Deng, Dongxu Li, Yunshen Wei, Baohong Lv, Junjie Yan, Lingpeng Kong, and Yiran Zhong. cosformer: Rethinking softmax in attention. In *International Conference on Learning Representations*, 2022.

[79] Nikita Kitaev, Lukasz Kaiser, and Anselm Levskaya. Reformer: The efficient transformer. In *International Conference on Learning Representations*, 2020.

[80] Krzysztof Choromanski, Valerii Likhosherstov, David Dohan, Xingyou Song, Andreea Gane, Tamas Sarlos, Peter Hawkins, Jared Davis, Afroz Mohiuddin, Lukasz Kaiser, et al. Rethinking attention with performers. *arXiv preprint arXiv:2009.14794*, 2020.

[81] Yu Zhang, Songlin Yang, Ruijie Zhu, Yue Zhang, Leyang Cui, Yiqiao Wang, Bolun Wang, Freda Shi, Bailin Wang, Wei Bi, Peng Zhou, and Guohong Fu. Gated slot attention for efficient linear-time sequence modeling. In *Proceedings of NeurIPS*, 2024.

[82] Yu Sun, Xinhao Li, Karan Dalal, Jiarui Xu, Arjun Vikram, Genghan Zhang, Yann Dubois, Xinlei Chen, Xiaolong Wang, Sanmi Koyejo, et al. Learning to (learn at test time): Rnns with expressive hidden states. *arXiv preprint arXiv:2407.04620*, 2024.

[83] Tianyuan Zhang, Sai Bi, Yicong Hong, Kai Zhang, Fujun Luan, Songlin Yang, Kalyan Sunkavalli, William T Freeman, and Hao Tan. Test-time training done right. *arXiv preprint arXiv:2505.23884*, 2025.

[84] Ali Behrouz, Peilin Zhong, and Vahab Mirrokni. Titans: Learning to memorize at test time. *arXiv preprint arXiv:2501.00663*, 2024.

[85] Han Guo, Songlin Yang, Tarushii Goel, Eric P Xing, Tri Dao, and Yoon Kim. Log-linear attention. *arXiv preprint arXiv:2506.04761*, 2025.

[86] Aaron Blakeman, Aarti Basant, Abhinav Khattar, Adithya Renduchintala, Akhiad Bercovich, Aleksander Ficek, Alexis Bjorlin, Ali Taghibakhshi, Amala Sanjay Deshmukh, Ameya Sunil Mahabaleshwarkar, et al. Nemotron-h: A family of accurate and efficient hybrid mamba-transformer models. *arXiv preprint arXiv:2504.03624*, 2025.

[87] Opher Lieber, Barak Lenz, Hofit Bata, Gal Cohen, Jhonathan Osin, Itay Dalmedigos, Erez Safahi, Shaked Meirom, Yonatan Belinkov, Shai Shalev-Shwartz, et al. Jamba: A hybrid transformer-mamba language model. *arXiv preprint arXiv:2403.19887*, 2024.

[88] Jusen Du, Weigao Sun, Disen Lan, Jiaxi Hu, and Yu Cheng. Mom: Linear sequence modeling with mixture-of-memories. *arXiv preprint arXiv:2502.13685*, 2025.

[89] Junxiong Wang, Daniele Paliotta, Avner May, Alexander Rush, and Tri Dao. The mamba in the llama: Distilling and accelerating hybrid models. *Advances in Neural Information Processing Systems*, 37:62432–62457, 2024.

[90] Michael Zhang, Simran Arora, Rahul Chalamala, Benjamin Frederick Spector, Alan Wu, Krithik Ramesh, Aaryan Singhal, and Christopher Re. Lolcats: On low-rank linearizing of large language models. In *The Thirteenth International Conference on Learning Representations*, 2024.

[91] Disen Lan, Weigao Sun, Jiaxi Hu, Jusen Du, and Yu Cheng. Liger: Linearizing large language models to gated recurrent structures. In *International conference on machine learning*, 2025.

[92] Jungo Kasai, Hao Peng, Yizhe Zhang, Dani Yogatama, Gabriel Ilharco, Nikolaos Pappas, Yi Mao, Weizhu Chen, and Noah A Smith. Finetuning pretrained transformers into rnns. In *Proceedings of the 2021 Conference on Empirical Methods in Natural Language Processing*, pages 10630–10643, 2021.

[93] Aviv Bick, Kevin Li, Eric P. Xing, J Zico Kolter, and Albert Gu. Transformers to SSMs: Distilling quadratic knowledge to subquadratic models. In *The Thirty-eighth Annual Conference on Neural Information Processing Systems*, 2024.

[94] Jean Mercat, Igor Vasiljevic, Sedrick Keh, Kushal Arora, Achal Dave, Adrien Gaidon, and Thomas Kollar. Linearizing large language models. *arXiv preprint arXiv:2405.06640*, 2024.

[95] Hanting Chen, Liu Zhicheng, Xutao Wang, Yuchuan Tian, and Yunhe Wang. Dijiang: Efficient large language models through compact kernelization. In *International Conference on Machine Learning*, pages 7103–7117. PMLR, 2024.

[96] Michael Zhang, Kush Bhatia, Hermann Kumbong, and Christopher Re. The hedgehog & the porcupine: Expressive linear attentions with softmax mimicry. In *The Twelfth International Conference on Learning Representations*, 2024.

[97] Barret Zoph and Quoc V Le. Neural architecture search with reinforcement learning. *arXiv preprint arXiv:1611.01578*, 2016.

[98] Han Cai, Tianyao Chen, Weinan Zhang, Yong Yu, and Jun Wang. Efficient architecture search by network transformation. In *Proceedings of the AAAI conference on artificial intelligence*, 2018.

[99] Xuan Shen, Pu Zhao, Yifan Gong, Zhenglun Kong, Zheng Zhan, Yushu Wu, Ming Lin, Chao Wu, Xue Lin, and Yanzhi Wang. Search for efficient large language models. *Advances in Neural Information Processing Systems*, 37:139294–139315, 2024.

[100] Youpeng Zhao, Ming Lin, Huadong Tang, Qiang Wu, and Jun Wang. Merino: Entropy-driven design for generative language models on iot devices. In *Proceedings of the AAAI Conference on Artificial Intelligence*, pages 22840–22848, 2025.

[101] Akhiad Bercovich, Tomer Ronen, Talor Abramovich, Nir Ailon, Nave Assaf, Mohammed Dabbah, Ido Galil, Amnon Geifman, Yonatan Geifman, Izhak Golan, et al. Puzzle: Distillation-based nas for inference-optimized llms. In *Forty-second International Conference on Machine Learning*, 2025.

[102] Pavlo Molchanov, Jimmy Hall, Hongxu Yin, Jan Kautz, Nicolo Fusi, and Arash Vahdat. Lana: latency aware network acceleration. In *European Conference on Computer Vision*, pages 137–156. Springer, 2022.

[103] Ruisi Cai, Saurav Muralidharan, Greg Heinrich, Hongxu Yin, Zhangyang Wang, Jan Kautz, and Pavlo Molchanov. Flextron: Many-in-one flexible large language model. *arXiv preprint arXiv:2406.10260*, 2024.

[104] Ruisi Cai, Saurav Muralidharan, Hongxu Yin, Zhangyang Wang, Jan Kautz, and Pavlo Molchanov. Llamaflex: Many-in-one llms via generalized pruning and weight sharing. In *The Thirteenth International Conference on Learning Representations*, 2025.

[105] Joshua Ainslie, James Lee-Thorp, Michiel de Jong, Yury Zemlyanskiy, Federico Lebron, and Sumit Sanghai. Gqa: Training generalized multi-query transformer models from multi-head checkpoints. In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing*, pages 4895–4901, 2023.

[106] Jingwei Zuo, Maksim Velikanov, Ilyas Chahed, Younes Belkada, Dhia Eddine Rhayem, Guillaume Kunsch, Hakim Hacid, Hamza Yous, Brahim Farhat, Ibrahim Khadraoui, Mugariya Farooq, Giulia Campesan, Ruxandra Cojocaru, Yasser Djilali, Shi Hu, Iheb Chaabane, Puneesh Khanna, Mohamed El Amine Seddik, Ngoc Dung Huynh, Phuc Le Khac, Leen AlQadi, Billel Mokeddem, Mohamed Chami, Abdalgader Abubaker, Mikhail Lubinets, Kacper Piskorski, and Slim Frikha. Falcon-h1: A family of hybrid-head language models redefining efficiency and performance. *arXiv preprint arXiv:2507.22448*, 2025.

A. Experimental Details

A.1. Final Model Architecture

The final Jet-Nemotron models are composed of a stack of blocks, each containing a Multi-Layer Perceptron (MLP) layer and an attention layer. The attention layer is selected from one of three types: full attention, sliding window attention, or JetBlock. The detailed architecture configurations are presented in Table 9.

	Jet-Nemotron-2B	Jet-Nemotron-4B
Total blocks	28	36
Full Attention Layers	No. 15, 20	No. 18, 21, 22, 28, 33
Sliding Window Attention Layers	No. 21, 22	No. 17, 20, 23, 24, 26
Vocabulary Size	151,643	151,643
Hidden Size	1,536	2,048
MLP Intermediate Size	8,960	11,008

Table 9 | The overall model architectures of Jet-Nemotron families.

The full attention and sliding window attention layers use grouped-query attention [105] and are configured as in Table 10. For sliding window attention layers, the window size is set to 1,152 in Jet-Nemotron-2B and 2,048 in Jet-Nemotron-4B.

Full Attention / SWA	Jet-Nemotron-2B	Jet-Nemotron-4B
Attention Head Number	12	16
Dimensions of Q/K/V	128	128
K/V Head Number	2	2
Position Embedding	RoPE	RoPE

Table 10 | The configurations of full-attention layers in Jet-Nemotron models.

The configuration of JetBlock are shown in Table 11 :

JetBlock	Jet-Nemotron-2B	Jet-Nemotron-4B
Q/K Dimension	96	128
V Dimension	256	256
Head Number	12	16
Convolution Kernel Size	4	4
DConv Generator Hidden Size	32	32

Table 11 | The configurations of JetBlock.

A.2. Experimental Costs

Table 12 summarizes the costs for PostNAS and training the Jet-Nemotron-2B model. We used 32 H100 GPUs in parallel. The reported GPU hours already account for the total number of devices.

A.3. Throughput Measurement

Throughout the experiments, we measure the maximum reachable prefilling and decoding throughput of Jet-Nemotron and the baselines on a single H100 GPU. This is achieved by adjusting the chunk size in chunk-prefilling [72] to maximize the decoding batch size without sacrificing the prefilling throughput. We list the optimized batch size and the corresponding chunk size for each model in Table 13. The prefilling context length is 64K. Since the KV cache memory dominates GPU usage during inference, by reducing the memory footprint per sequence, smaller caches allow more sequences to be processed in parallel, greatly boosting generation throughput.

		Tokens (B)	ZFLOPs	Time (H100 GPU Hours)
PostNAS	Full Attention Placement and Elimination	50	0.8	808
	Linear Attention Block Selection	50	4.0	3120
	New Attention Block Design	50	0.8	624
	Harware-Aware Arch Search	50	7.2	5616
Training	Stage1	50	0.8	624
	Stage2	350	5.6	7536

Table 12 | Experimental Costs for PostNAS and training the Jet-Nemotron-2B model.

Model	Batch Size	Chunk Size
Qwen2.5-1.5B	32	8,192
Qwen3-1.7B	8	16,384
Llama3.2-1B	32	4,096
MiniCPM-2B-128K	2	2,048
Pythia-2.8B	2	16,384
Smollm2-1.7B	4	16,384
Mamba2-2.7B	128	1,024
RWKV7-1.5B	256	2,048
Rec.Gemma-2B	128	512
Gemma3n-E2B	64	4,096
Gemma2-2.6B	16	2,048
Hymba-1.5B	64	512
Zamba2-1.2B	8	8,192
Jet-Nemotron-2B	128	2,048
Jet-Nemotron-4B	64	1,024

Table 13 | **Hyper-Parameters in Efficiency Measurement.** We adjust the chunk size to maximize decoding batch size without compromising prefilling throughput.

B. Additional Results

B.1. Controlled Study on Training Data

To exclude the influence of training data, we continually pre-train the baseline models (Qwen2.5, RWKV-7, and Mamba-2) on Jet-Nemotron 's training dataset to provide a more comprehensive evaluation. The results in Table 14 show that Jet-Nemotron-2B outperforms all these finetuned baseline models by a significant margin.

Model	MMLU	Math	Commonsense	Retrieval
Qwem2.5-1.5B-continual	56.7	37.6	59.8	71.5
Mamba2-2.7B-continual	41.0	22.5	56.9	55.9
RWKV7-1.5B-continual	49.8	25.2	59.3	57.2
Jet-Nemotron-2B	**59.6**	**40.2**	**61.7**	**73.6**

Table 14 | **Controlled Study on Training Data**. All models are pre-trained or continually pre-trained on the Jet-Nemotron stage-2 training corpus discussed in Section 3.1.

B.2. Throughput Results on Lower-End Hardware

We measure the throughput of Jet-Nemotron-2B and Qwen2.5-1.5B on the NVIDIA Jetson Orin (32GB) and NVIDIA RTX 3090 GPUs with a context length of 64K. Results in Table 15 show that Jet-Nemotron-2B achieves 8.84× and 6.50× speedups over Qwen2.5-1.5B on the Jetson Orin and RTX 3090 GPUs, respectively.

Hardware	Qwen2.5-1.5B (Tokens/s)	Jet-Nemotron-2B (Tokens/s)	SpeedUp
Orin	6.22	55.00	8.84
3090	105.18	684.01	6.50

Table 15 | Throughput Results on Jetson Orin (32GB) and NVIDIA RTX 3090 GPUs.

Model	Throughput	Accuracy ↑					
	(token/s) ↑	MMLU	MATH	Common.	Retrieval	Code	Long-Context
Falcon-H1-1.5B [106]	223	60.5	40.1	59.9	73.5	56.0	40.7
Falcon-H1-1.5B-deep [106]	66	<u>63.5</u>	46.8	60.6	<u>74.6</u>	<u>60.3</u>	33.4
Jet-Nemotron-2B	2,885	60.8	<u>49.6</u>	<u>62.0</u>	74.2	59.5	<u>41.1</u>
Jet-Nemotron-4B	1,271	**65.2**	**51.3**	**64.7**	**76.2**	**63.5**	**43.9**

Table 16 | **Comparison with Falcon-H1.**

B.3. Comparison to Falcon-H1

We compare our work with the concurrent Falcon-H1 [106], a hybrid model that incorporates Mamba2 [50] and full attention. Unlike Jet-Nemotron, which alternates between component types at the layer level, Falcon-H1 employs a head-wise hybrid strategy. As shown in Table 16, Jet-Nemotron-2B outperforms Falcon-H1-1.5B and is comparable to Falcon-H1-1.5B-deep in accuracy, while achieving significantly higher generation throughput. Jet-Nemotron-4B outperforms both the two Falcon-H1 models while still achieves higher generation throughput. This efficiency gap arises because the head-wise strategy requires sequential computation of Mamba2 and full attention operations within a single layer, thereby limiting parallelism. The "-deep" variant further exacerbates this issue by reducing model width in favor of greater depth.